

Assignment No.2
Name : Shivani Gaikwad
Roll No. 43315
Batch : P-11

1. Import the necessary packages

```
In [1]: from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import mnist
from tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np
import argparse
```

2. Grab the MNIST Dataset

```
In [2]: print("Accessing mnist...")
((trainX, trainY), (testX, testY)) = mnist.load_data()

# Flatten the images
trainX = trainX.reshape((trainX.shape[0], 28 * 28 * 1))
testX = testX.reshape((testX.shape[0], 28 * 28 * 1))

# Scale data to the range of [0, 1]
trainX = trainX.astype("float32") / 255.0
testX = testX.astype("float32") / 255.0

# convert labels from integers to vectors
le = LabelBinarizer()
trainY = le.fit_transform(trainY)
testY = le.fit_transform(testY)
```

Accessing mnist...

3. Feed Forward Neural Network

```
In [3]: model = Sequential()
model.add(Dense(256, input_shape = (784,), activation = "sigmoid"))
model.add(Dense(128, activation = "sigmoid"))
model.add(Dense(10, activation = "softmax"))
```

4. Adding SGD Optimizer and training

```
In [4]: # Adding optimizer
# Learning rate --> 0.01
sgd = SGD(0.01)

# Fitting and training
model.compile(loss="categorical_crossentropy", optimizer=sgd, metrics=["accuracy"])
H = model.fit(trainX, trainY, validation_data=(testX, testY), epochs=100, batch_size=32)

Epoch 57/100
469/469 [=====] - 2s 5ms/step - loss: 0.3169 - accuracy: 0.9090 - val_loss: 0.3057 - val_accuracy: 0.9115
```

```
In [5]: print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	200960
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 10)	1290

=====
 Total params: 235,146
 Trainable params: 235,146
 Non-trainable params: 0

None

5. Evaluating the model

```
In [6]: predictions = model.predict(testX, batch_size=128)
print(classification_report(testY.argmax(axis=1), predictions.argmax(axis=1), tar
```

	precision	recall	f1-score	support
0	0.94	0.98	0.96	980
1	0.97	0.98	0.97	1135
2	0.92	0.91	0.91	1032
3	0.90	0.91	0.91	1010
4	0.92	0.93	0.92	982
5	0.91	0.86	0.88	892
6	0.93	0.94	0.94	958
7	0.93	0.92	0.93	1028
8	0.89	0.90	0.89	974
9	0.91	0.90	0.91	1009
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

6. Plotting training loss and accuracy

```
In [7]: plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
```

Out[7]: <matplotlib.legend.Legend at 0x1966024e850>

